The interactive quiz supports 20 questions per round, so this assessm
1–20) and Round 2 (Questions 21–30). Complete Round 1 and use the next-quiz option to continue.

14 of 20

A trainee wants to test prompts against supported generative AI models without writing an application. Which Databricks interface is designed for this?

A Databricks SQL editor

B MLflow experiment tracking

C Lakeflow pipeline editor

D AI Playground

Next

Ask ChatGPT

Medium

Search

ENG
IN

he interactive quiz supports 2[illegible]

–20) and Round 2 (Questions 21–30). Complete Round 1 and use the next-quiz option to continue.

In Databricks generative AI features, what does LLM stand for? 13 of 20

A Lakehouse Logic Model

B Language Learning Mechanism

C Linked Language Metadata

D Large Language Model

Next

1–20) and Round 2 (Questions 21–30). Complete Round 1 and

12 of 20

A team shares approved datasets across Databricks workspaces. Which capability is provided by Unity Catalog?

A Automatic removal of all cloud storage policies

B Centralized permissions, discovery, and lineage

C Automatic provisioning of compute for every query

D Automatic replacement of every existing Delta table

Next

A learner evaluates an ML model and wants to record its accuracy for a specific training run. What should be logged in MLflow?

A A run metric

Accuracy is a numeric evaluation measurement and can be logged as an MLflow run metric.

B An experiment description

C A model alias

D A registered model version

Next

A learner evaluates an ML model and wants to record its accuracy for a specific training run. What should be logged in MLflow?

A. A run metric

B. An experiment description

C. A model alias

D. A registered model version

Next

A Lakeflow Spark Declarative Pipeline should identify orders whose order_id is missing. Which feature is designed to define and measure this data-quality condition?

A Streaming checkpoint location

B Pipeline event-log retention

C Data-quality expectation

D Pipeline full-refresh setting

Next

Two updates arrive for the same customer in the wrong order. Which source information is best for applying the changes in the intended sequence?

8 of 20

A The cloud file's upload folder name

B The processing cluster's start time

C A trusted event sequence or ordering timestamp

D The target table's creation date

Next

Ask ChatGPT

Medium

Search

The interactive

1–20) and Round 2 (Questions 21–30). Complete Round

7 of 20

A CDC source sends insert, update, and delete events. Why does the pipeline need the operation type for each event?

A To determine how the corresponding target record should change

B To define the table retention duration

C To choose the Spark runtime for a pipeline

D To select a partition size for the target table

Next

6 of 20

A Delta write includes an incompatible column type. Which Delta Lake capability normally rejects this unexpected schema change?

A Schema enforcement

B Automatic file compaction

C Schema evolution

D Delta time travel

Next

The interactive quiz supports 20 questions per round.
1–20) and Round 2 (Questions 21–30). Complete Round 1 and use the next-quiz option to continue.

5 of 20

Before investigating an incorrect Delta update, an engineer wants a list of previous table operations and versions. Which command should be used?

A DESCRIBE HISTORY

B SHOW CREATE TABLE

C SHOW TBLPROPERTIES

D DESCRIBE DETAIL

Next

Which Delta write mode adds new rows while keeping existing rows, provided the input schema is compatible?

4 of 20

A Append

B Ignore

C Overwrite

D Error if exists

Next

The interactive quiz supports 20 questions per [illegible]
1–20) and Round 2 (Questions 21–30). Complete Round 1 and use the next-quiz option to continue.

Delta tables store data files in Parquet format. What allows Delta Lake to track committed changes and support reliable transactions?

3 of 20

A Unity Catalog lineage

B Delta data-file metadata

C Delta table statistics

D Delta transaction log

Next

Ask ChatGPT

Medium

Search

SAMSUNG

The interactive quiz supports 20 questions per round, so this assessment
1–20) and Round 2 (Questions 21–30). Complete Round 1 and use the next-quiz option to continue.

Which Lakeflow Jobs feature lets a workflow use different input dates without changing notebook code?

2 of 20

A Job parameters

B Job cluster policies

C Job scheduling

D Task dependency settings

Next

-20) and

What is the purpose of a trigger in a Databricks Lakeflow Job?

1 of 20

A Set the number of worker nodes for each task

B Determine when a job should start

C Control the sequence of tasks after a run starts

D Specify how many times a failed task retries

Next

Touch panel re
IPv4: 10.201.230
IPv6:

An insurance analytics Delta table was updated incorrectly at 10 AM. At 10:30 AM, a second job committed legitimate corrections to other rows. The trainee is asked to analyze possible recovery paths, not to modify Production. The Delta table has history available for both operations. The manager wants the 10:30 AM corrections retained. A complete table restore to 9 AM would also remove those later corrections from the current state. The team requires a documented comparison before a change request is approved. Which proposal meets the stated recovery objective?

A Restore to the earlier version immediately because newer commits can be reapplied automatically

B Delete the Delta log version associated with the error to undo only that update

C Compare versions, identify the bad changes, and propose targeted repair preserving later valid updates

D Run VACUUM to reverse the erroneous operation while retaining every later commit

+ Ask ChatGPT Medium

Search

A retail inventory Delta table receives frequent small writes throughout the day. There are thousands of active small Parquet files under the table. The morning dashboard repeatedly scans the same date range and has become slow. A new engineer recommends immediately running VACUUM to compact the table. The team must preserve transaction history according to its retention policy. An authorized owner wants a maintenance suggestion that improves file layout without treating old files as the root cause. The engineer may evaluate the proposal in Dev before any production change. Which recommendation is most suitable?

13 of 14

Touch panel required
IPv4: 10.201.230.219
IPv6:

A Run RESTORE every morning because restoring reduces active file counts

B Assess OPTIMIZE to compact active small files and verify performance before an approved change

C Run VACUUM because it combines active small files into larger ones

D Disable Delta statistics collection so file scan become faster

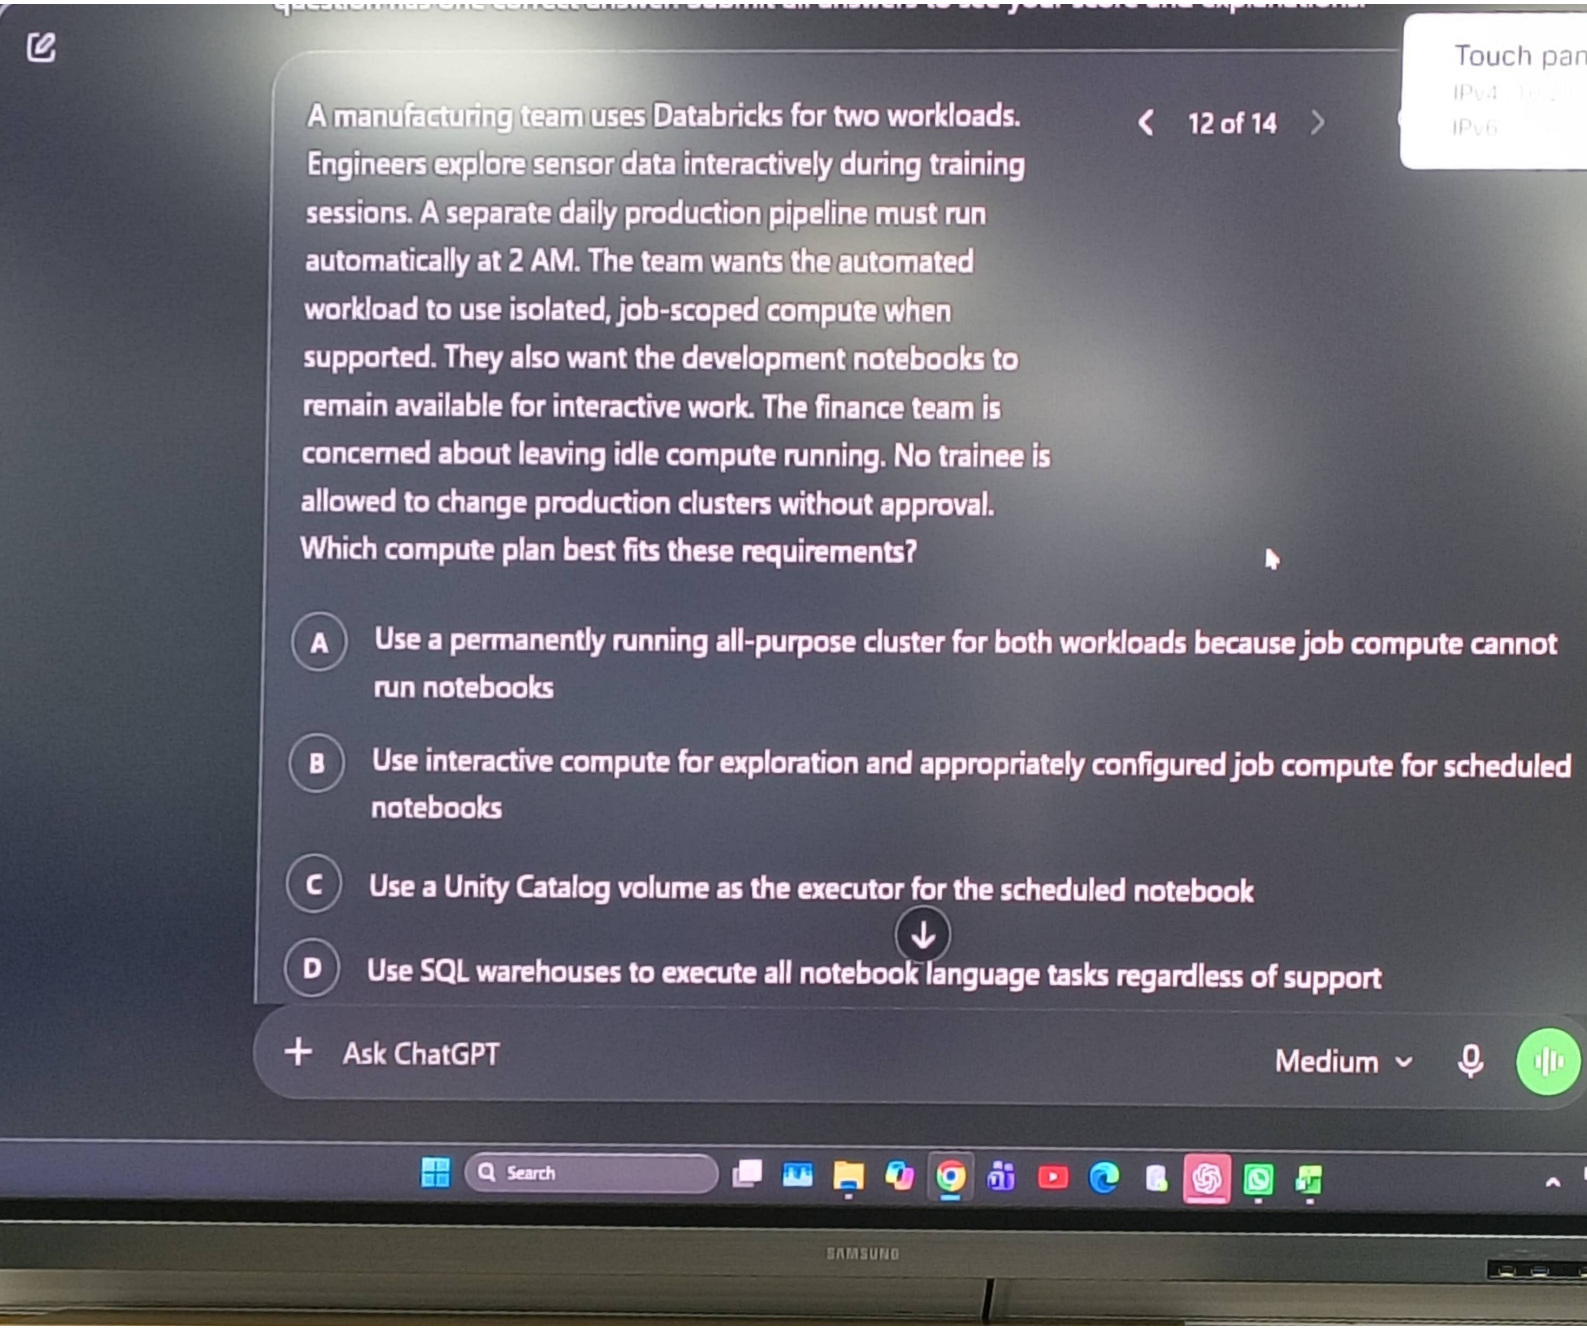
Ask ChatGPT

Medium

Search

ENG
IN

1524


Touch pan
12 of 14
A manufacturing team uses Databricks for two workloads. Engineers explore sensor data interactively during training sessions. A separate daily production pipeline must run automatically at 2 AM. The team wants the automated workload to use isolated, job-scoped compute when supported. They also want the development notebooks to remain available for interactive work. The finance team is concerned about leaving idle compute running. No trainee is allowed to change production clusters without approval. Which compute plan best fits these requirements?
A Use a permanently running all-purpose cluster for both workloads because job compute cannot run notebooks
B Use interactive compute for exploration and appropriately configured job compute for scheduled notebooks
C Use a Unity Catalog volume as the executor for the scheduled notebook
D Use SQL warehouses to execute all notebook language tasks regardless of support
Ask ChatGPT
Medium
Search
SAMSUNG

A delivery platform runs an interactive Databricks notebook during development. The notebook reads a governed Delta table, groups shipment events, and produces a report. Its Spark driver starts the application and coordinates work. Several executors process partitions of shipment data. The trainee notices that the job becomes slow when a single partition is much larger than the others. The cluster has spare workers, but one executor remains busy for much longer. A senior engineer asks the trainee to explain the architecture and likely bottleneck. Which explanation is the most technically accurate?

A Driver coordinates the application; executors process tasks; skew can leave one task slow

B Driver processes every input row while executors only store cached output

C The cluster manager guarantees equal-duration tasks despite unequal data sizes

D Executors schedule their own independent jobs without driver coordination

question has one correct answer. Submit all answers to see your score and explanations.

Touch panel
IPv4: 10.201.23
IPv6:

10 of 14

A scheduled Databricks workflow fails while loading yesterday's records. The previous task may already have written part of the batch, and the assigned fresher has monitoring access but no permission to restart production jobs. What is the safest first response?

A Inspect run and batch state with permitted access, then escalate for an authorized retry decision

B Reset the streaming checkpoint before examining the failure

C Rerun the entire job immediately without checking already-written data

D Modify the source batch identifier so the same input appears new

Next

A project stores a product CSV in a Unity Catalog external volume. A new analyst can use the catalog and schema but receives an access-denied error reading the volume file. The analyst needs read access only and must not administer the volume. Which access change should the authorized owner evaluate?

A Request READ VOLUME while retaining the required USE privileges on parent objects

B Request CREATE VOLUME on the schema instead of file read access

C Request ownership of the external location for all analysts

D Request WRITE VOLUME so files can be read but not changed

Next

IPv6:

8 of 14

An inventory Delta table received an incorrect update, but several valid transactions were committed afterward. A trainee has read access and is asked to help the owner investigate. Which response best preserves those newer valid transactions?

A Restore the whole table to the version before the error without reviewing later commits

B Use VACUUM as a rollback command to reverse the bad update

C Inspect Delta history, isolate erroneous changes, and propose an approved targeted fix

D Delete older data files manually, keeping only the most recent version

Next

column, and

schema. Which write approach best fits this planned

replacement?

A Use an approved overwrite operation with the schema-overwrite option enabled

Overwrite with overwriteSchema can replace the dataset and intentionally update its schema, subject to applicable table restrictions.

B Use append with mergeSchema, leaving old rows in place

C Use a Delta UPDATE statement to replace the full table and add columns

D Use overwrite without a schema-change option and assume the new column is accepted

Next

Touch panel
IPv4: 10.201.23
IPv6:

An equipment company receives changes for machine_id M12 in this arrival order: update sequence 105, then update sequence 103. The source sequence reflects the true order of business events, but network delays can change their arrival order. The target should represent the newest state without allowing older events to overwrite it. Which mechanism is most appropriate?

A Apply each source row according to its alphabetical status text

B Apply CDC updates in whichever order the tasks finish processing

C Use the trusted source sequence to order CDC updates for each key

D Use the object-storage arrival timestamp to decide the newest business event

Next

Ask ChatGPT

Medium

Here is your 14-question

question has one correct answer. Submit all answers to see your

IPv6:

7 of 14

A training project stores inspection records in a Delta table. A full refresh should replace the existing rows with a newly validated batch. The new batch has an additional approved column, and the team needs to deliberately update the target schema. Which write approach best fits this planned full replacement?

A Use an approved overwrite operation with the schema-overwrite option enabled

B Use append with mergeSchema, leaving old rows in place

C Use a Delta UPDATE statement to replace the full table and add columns

D Use overwrite without a schema-change option and assume the new column is accepted

Next

Medium

6 of 14

An equipment company receives changes for machine_id M12 in this arrival order: update sequence 105, then update sequence 103. The source sequence reflects the true order of business events, but network delays can change their arrival order. The target should represent the newest state without allowing older events to overwrite it. Which mechanism is most appropriate?

A Apply each source row according to its alphabetical status text

B Apply CDC updates in whichever order the tasks finish processing

C **Use the trusted source sequence to order CDC updates for each key**

Sequencing CDC events by a trusted business-ordering field prevents late-arriving older updates from replacing newer state.

D Use the object-storage arrival timestamp to decide the newest business event

Ask ChatGPT

Medium

IPv6:

Which Databricks feature is designed to display charts and business metrics from query results?

4 of 14

A SQL query history

B SQL query profile

C **AI/BI dashboard**

AI/BI dashboards visualize query results, metrics, and business indicators.

D SQL alert

Next

IPv6:

A trainee runs a SELECT query in the Databricks SQL editor. 3 of 14

Which compute resource executes this query?

A SQL warehouse

Databricks SQL uses a SQL warehouse to execute queries submitted from the SQL editor.

B All-purpose interactive cluster

C Job compute used by a scheduled notebook task

D Model serving endpoint

Next

...answers to see your score and explanations.

Which Unity Catalog access practice best demonstrates least privilege for a new trainee?

2 of 14

A **Grant only the required permissions on relevant data objects**
Least privilege grants only the access necessary for assigned work.

B Give the trainee ownership of the full catalog

C Grant write privileges whenever the trainee only needs to read

D Provide broad privileges and remove them after the project

Next